

GW — lab notebook (gravitational-wave program)

K. Bhatia ORCID 0009-0007-4447-6325 repo: GW (ported from kunalb541/cosmo2, E-numbering preserved)
Data of record: GWTC-2.1/3 PE (GWOSC/Zenodo, 76 events); GWTC-4.0 PE (Zenodo 16053484, 86 events);
LVK TGR products (Zenodo 17461225: imr/par/rin/liv); GWTC-3 cosmology icarogw (Zenodo 5645777); PTA
chains: NANOGrav 15yr fig-1 release, EPTA DR2 (Zenodo 8091568), PPTA DR3 (github). Every battery locked-
preregistered before running; every headline independently re-derived; numbers of record in **results/**, per-battery
detail in **reports/**.

Black-hole posterior geometry: the event atlas and its laws (E38, E40–E44)

E38 GW posterior-geometry atlas. Value/shape geometry on real compact-binary posteriors: all 76 confident GWTC-1/2.1/3 events (1.54M samples; GWOSC/Zenodo PE releases), 6 black-hole planes each — (m_1, m_2) , $(\mathcal{M}, q)$, (q, χ_{eff}) , (M_f, a_f) , $(d_L, \cos \iota)$, (z, M_{tot}) ; 74/76 produce all six. The foundation for E40–E47.

E40 GW chirp-mass lawfulness. Is the (m_1, m_2) posterior degeneracy just the constant-chirp-mass line? Predict each event's orientation from the constant- $\mathcal{M}$ tangent at its median masses, vs a constant-total-mass null, 75 events. PASS: median $|\Delta\psi_{\text{chirp}}| = 7.49^\circ$ vs total-mass 20.38° ; chirp wins 61/75; on the elongated subset (axis ratio ≥ 3 , $n=28$) 4.68° ; Spearman(error, axis-ratio) = -0.42 (more elongated $\rightarrow$ tighter, as predicted). A rare POSITIVE geometric-lawfulness result — geometry is lawful where a dominant physical combination (chirp mass) exists.

E41 GW plane lawfulness map. A-priori correlation-sign predictions for all 6 planes, derived by 12 physicist/verifier agents blind to the data, tested against measured within-event ρ (75 events). 3/6 HIT:

E42 BH population + precision lawfulness. Precision test: chirp-mass fractional precision $\sigma_{\mathcal{M}}/\mathcal{M}$ vs SNR (75 events). PASS: $r = -0.50$ ($p=7\times 10^{-7}$), OLS slope -1.22 (Fisher -1), and partial correlation with $\log \mathcal{M}$ at fixed SNR = $+0.94$ (the cycle-count law $N \propto \mathcal{M}^{-5/3}$). Posterior SIZE is lawful — completing orientation (E40) + correlation (E41) + size (E42). Population map: χ_{eff} mean $+0.076$, 36% median-negative, m_1 peaks near 10 and $35 M_\odot$, 7 IMBH-scale remnants, 1 BNS; the geometric outliers are the physically extreme events (GW190814 $q=0.11$, GW191219 $q=0.04$ NSBH, GW200115 NSBH, GW190425 BNS).

E43 χ_{eff} per-event negatives. Does any event securely have $\chi_{\text{eff}} < 0$ (dynamical-formation signature)? From the full posteriors (75 events): 0 events with $P(\chi_{\text{eff}} < 0) > 0.90$ (GW191109 sits at exactly 0.900, 90% CI $[-0.60, +0.12]$ spans 0), 0 with a 95% CI below 0. NULL: E42's 36% median-negative is largely measurement scatter, not per-event-secure negatives.

E44 χ_{eff} hierarchical. Does the POPULATION require support below 0? Gaussian fit (flat and induced-prior-reweighted, 75 events) + a positive-truncation comparison. $\mu \approx 0.06\text{--}0.09$, $\sigma \approx 0.045\text{--}0.061$, negative fraction median $\sim 9\%$ — but a $\chi_{\text{eff}} \geq 0$ truncated population fits as well or better ($\Delta \ln L = -1.0/-0.74$), so the negative tail is a Gaussian-tail artifact, NOT required. Reinforces E43. (Does not overturn the LVK negative tail: a simplified pipeline — Gaussian-only, approximate prior, no selection function — simply does not independently demand it.)

Testing general relativity with black holes (E45–E47, E55, E57–E60)

E45 IMR-consistency test. LVK GWTC-3 Tests-of-GR: the inspiral- vs post-inspiral-inferred final mass/spin, deviation $(\Delta M_f/M_f, \Delta a_f/a_f)$, GR = (0, 0), 18 events. GR UPHELD — D_1 combined deviation $(-0.035, +0.002)$, GR at credible level 0.671; D_2 17/18 consistent (outlier GW190814 has an absent post-inspiral, a known systematic). D_3 (the framework's added test): the mean-deviation directions are COHERENT (Rayleigh $p = 0.001$, 13/18 with positive final-spin deviation), only partly along the shared degeneracy — a coherent sub-threshold waveform systematic that per-event tests miss.

E46 parameterized PN deviations. Fractional PN inspiral-phase deviations $\delta\hat{\varphi}_i$ (GR = 0), 9 events $\times$ 10 coefficients (SEOBNRv4). GR NOT violated. The cross-event sign-coherence test isolates **dchi0** (leading 0PN):

9/9 events positive, Bonferroni $p = 0.004$, combined $+0.109 [0.040, 0.177]$ excluding 0 — the documented low-PN WAVEFORM SYSTEMATIC of a single-waveform analysis, not new physics. The non-coherent combined exclusions (`dchi3/51/61`, only 2–4/9 positive) are naive-product artifacts the coherence test correctly discounts. Single waveform; multi-waveform robustness (not run) would confirm the systematic origin.

E47 ringdown no-hair test. Black-hole spectroscopy: fractional deviation of the 221 overtone frequency/damping ($\delta f, \delta \tau$) from the Kerr prediction, GR = (0, 0), 22 events. GR/Kerr UPHELD — D_1 combined 0.75σ from GR ($Q = 0.246$, via a robust Gaussian product); D_2 22/22 consistent; D_3 ISOTROPIC (Rayleigh $p = 0.79$), no coherent deviation (unlike E45/E46 — the coherent systematics are inspiral-specific, absent in the overtone). A first-pass histogram-density product spuriously reported a 99.9% “no-hair violation” that contradicted D_2/D_3 (an event at $Q=0.987$ was only 0.25σ from GR); the Mahalanobis cross-check caught the binning artifact and the Gaussian product corrects it.

E55 multi-waveform robustness. E40’s chirp-mass orientation across two independent waveform families (IMRPhenomXPHM vs SEOBNRv4PHM, 59 events): PASS — the law holds for both (median $|\Delta\psi_{\text{chirp}}| 9.1^\circ/6.1^\circ$) and the orientations agree to 3.3° (elongated). Real geometry is waveform-robust; E46’s low-PN `dchi0` coherence, by contrast, is a known waveform systematic (would not survive). DATA-BLOCK (honest): a direct same-event multi-waveform test of the E45/E46 GR-test coherences is impossible with GWTC-3 (the parameterized set is FTI/SEOBNR-only; TIGER/IMRPhenom was delayed) — would need the GWTC-2 (different events).

E57 black-hole area theorem. Hawking’s 2nd law from GWTC-3: $A_{\text{final}} \geq A_{\text{initial}}$ using the inspiral progenitors (initial areas) vs the ringdown remnant (final area), disjoint/ independent (Isi 2021 setup). Only 4 loud events carry a ringdown remnant. HELD: no violations, area increases 5–79%, GW170104 strong single-event confirmation 0.92, stacked 0.92σ . Honest: crude independent-pairing + broad inspiral spin priors give GW150914 $P=0.74$, below Isi’s 0.97 — a consistency check, not a stringent bound.

E58 massive graviton / Lorentz-violation. GW-propagation gravity test: modified dispersion $E^2 = p^2 c^2 + A p^\alpha$ over 12 events $\times$ 8 orders ($\alpha=0$ = graviton mass). GR UPHELD: no Lorentz-violation/dispersion signal in any event or order. Graviton mass (per-event, robust): best single-event $m_g < 6.7 \times 10^{-24}$ eV/ c^2 , median 1.8×10^{-23} (LVK combined 1.27×10^{-23}). Honesty catch: a naive posterior-product combined bound gave an absurd 5.9×10^{-27} (3 orders too tight, grid-unstable) — LVK use a flat-in- A prior and the prior samples are empty in the release, so it cannot be divided out; the per-event bounds are reported and the LVK combined quoted, not a fabricated number.

E59 ringdown mode consistency. Does adding the 221 overtone shift the remnant (M_f, a_f) from the 220-fundamental value? `rin.zip` Kerr_220 vs plain Kerr_221, 22 events. A naive test “fails” (7 events $> 2\sigma$, max 6.1), BUT the shift is COHERENT — M_f decreases in 22/22 (100%) and a_f in 21/22 (95%) with the overtone. A uniform all-event down-shift is the known ringdown- overtone SYSTEMATIC (start-time/modeling; the GW150914 overtone controversy), NOT a no-hair violation. The coherence lens (as in E45/E46) prevents a false violation claim; the clean no-hair result remains E47 (221 deviation params ~ 0 , isotropic).

E60 model-independent ringdown no-hair. The strongest-in-principle ringdown test: the MEASURED 220 QNM frequency (damped-sinusoid fit, no Kerr assumption) vs the Kerr prediction $f_{220}(M_f, a_f)$ from the INSPIRAL remnant (Berti fit) — independent data. Kerr CONSISTENT: no event deviates $> 2\sigma$ (max $|z|=0.85$), combined 0.34σ . GW150914 clean: Kerr 236.6 vs measured 246.8 Hz (+5%). Weak (only 3 events have both a DS ringdown and an inspiral remnant; GW170814’s DS frequency is broad/poorly constrained), GW150914 carries the confirmation. Ringdown no-hair trilogy: E47 (deviation params ~ 0) + E59 (remnant shift = systematic) + E60 (frequency matches Kerr) all uphold the Kerr nature.

Deriving the mass-plane law; the curved chirp-mass law (E65, E67)

E65 the E40 law derived — locked Fisher-rotation REFUTED; the posterior IS the constant- $\mathcal{M}_c$ curve. Locked hypothesis: E40’s residual pattern (heavy events fail toward constant- M_{tot}) is a Fisher-mixture rotation — inspiral cycles measure $\mathcal{M}_c$, the merger band measures M_{tot} , so orientation should rotate with detector-frame mass. Model: two-direction Fisher mixture, analytic aLIGO PSD, one constant κ fit ONLY on the 14 locked E40-prereg events, scored out-of-sample on the other 61. OUTCOME: D1 marginal pass (7.45 vs

7.72°), D2 FAIL, D3 (parameter-free sign test) FAIL decisively — measured ψ does NOT sit between the chirp and total tangents; rotation fractions are mostly NEGATIVE (light elongated events overshoot PAST the chirp tangent, away from total), a sign no rotation model can produce. What the failure revealed (post-hoc, labelled): the posterior lies along the constant- $\mathcal{M}_c$ CURVE; the principal axis of a curved arc is its chord, not the median-point tangent. Zero-free-parameter curved law (uses only the $\mathcal{M}_c$ median + q marginal): median error 4.85 vs 7.49° (all 75), **0.81** vs 4.68° (elongated, 6×); GW190728’s 14° “violation” reproduced to 1°. On-record GWTC-4 out-of-sample prediction: $< 2^\circ$ for elongated O4a events. 5 contract tests; transient ernno-60 read handled by retry; headlines re-derived by hand.

E67 GWTC-4 out-of-sample: the curved chirp-mass law PASSES at 1.26°. The E65 prediction — locked on record (commit 8d09cbf) from GWTC-1/2.1/3 data only, prereg pushed BEFORE any GWTC-4 file was opened — scored on the fully disjoint GWTC-4.0 O4a catalog (Zenodo 16053484; 86 events, all processed, zero dropped, all files h5py-validated). **D1 PASS: median $|\psi_{\text{curve}} - \psi_{\text{meas}}| = 1.26^\circ$ on the 19 elongated events (threshold 2°)** — 18 of 19 land at 0.2–3.3° while the tangent carries systematic 4–11° errors that the curvature term removes (GW231118_090602: 10.7 $\rightarrow$ 0.5; GW231020: 8.0 $\rightarrow$ 0.2; GW230529 mass-gap: 2.9 $\rightarrow$ 0.3). D2 PASS: curve beats tangent on all 86 (7.53 vs 9.36°). D3: the E40 tangent law itself replicates out-of-sample (86% wins vs 81%; Spearman -0.40 vs -0.42). The E65 finding is hereby PROMOTED from post-hoc to out-of-sample confirmed: the (m_1, m_2) posterior geometry of GW events is the constant- $\mathcal{M}_c$ curve (zero free parameters — each prediction uses only the event’s own median source $\mathcal{M}_c$ and q marginal). Headlines re-derived independently + one event recomputed end-to-end by hand (exact match).

Gravitational-wave cosmology: the spectral-siren lever (E66)

E66 spectral-siren H_0 obstruction — REFUTED at GWTC-3 precision. Radical claim (E61’s obstruction logic on GW): the spectral-siren H_0 program is blocked by the $\sim 33 M_\odot$ mass-peak location μ_g acting as an uncalibrated ruler. Test: linear-response couplings on the REAL LVK GWTC-3 joint cosmology+population posterior (icarogw release, $n = 13,533$; power-law+ peak, restricted flat- Λ CDM). Locked stop-and-diagnose fired mid-run: the multivariate response $C_{AB}C_{BB}^{-1}$ (partial) disagreed with the mandatory decile cross-check (marginal) — an apples-to-oranges comparison in the implementation, caught exactly as the two-estimator rule intends; fixed to the marginal degeneracy track, after which all sanity gates pass. RESULT: a $4 M_\odot$ coherent μ_g drift moves H_0 by only -0.51 km/s/Mpc ($< 1 \Rightarrow$ REFUTED; robust $-0.41/-0.32$ across SNR cuts) — an order of magnitude below the CMB-vs-SH0ES gap and inside the chain’s own ± 3.3 statistical error. What survives: μ_g IS the single largest systematic lever (mass group 0.38 vs rate group 0.22 km/s/Mpc per 1σ), so the obstruction becomes live when spectral-siren precision reaches ~ 1 km/s/Mpc (O5/XG) — then an independently pinned or z -dependent peak is mandatory. Conservative caveat: the restricted $H_0 \in [65, 77]$ prior compresses couplings (lower bounds), but the verdict does not hinge on it.

The nanohertz band: the GW background across pulsar-timing arrays (E68)

E68 cross-PTA GWB anatomy — pure degeneracy-sliding; γ marches with time span. First value/shape battery in the nanohertz band, on real validated chains: NANOGrav 15yr (HD-correlated), EPTA DR2full (HD), PPTA DR3 (CRN), + EPTA DR2new as robustness; CPTA has no public GWB chain. D1: all pairs mutually consistent (1.0–1.8 σ), Gaussian-faithful (sliced-W ratios 0.98–1.01). D2 HEADLINE: the constraint geometry is UNIVERSAL ($\rho(A, \gamma) = -0.85..-0.95$, $\psi = 110\text{--}116^\circ$ — one measurement kernel) and the pairwise mean offsets lie **94.6–99.9%** along the common A - γ degeneracy axis (mean 0.98; locked threshold 2/3) $\Rightarrow$ classified degeneracy-sliding: everything the PTAs disagree about lives along the direction a common power-law fit slides when read over different bands. D3: the SMBHB circular prediction $\gamma = 13/3$ is excluded $> 2\sigma$ by NANOGrav ALONE ($z = +3.03$; EPTA-full +0.29, PPTA +1.25) — the SH0ES pattern (E52): one experiment carries the narrative; no combined number fabricated (E58 lesson). POST-HOC (labelled): $\hat{\gamma}$ is STRICTLY MONOTONIC in observing span (10.3 $\rightarrow$ 24.7 yr: 2.79 $\rightarrow$ 3.25 $\rightarrow$ 3.87 $\rightarrow$ 4.21), including EPTA’s internal DR2new-vs-DR2full split (1.89 σ , same telescopes, offset also along the degeneracy) — the coherent signature of a spectrum that flattens at high frequency relative to a single power law (or a span-correlated systematic). Successor prediction on record: longer-span subsets of the same array should fit systematically steeper γ (preregisterable

at IPTA DR3 / NG 20yr). 4 contract tests; every headline independently re-derived; all probe values match the publications.

GW data beyond black holes: the neutron-star EOS transfer (E63)

E63 neutron-star mass-radius geometry. First transfer of the value/shape framework out of cosmology, into the neutron-star equation of state, in the (M, R) plane. Real posteriors: NICER pulse-profile for the low-mass star J0030 (Riley 2019 X-PSI ST+PST; Miller 2019 2-oval/3-oval) and the high-mass star J0740 (Riley 2021; Miller 2021), plus GW170817 tidal Λ mapped to (M, R) via the quasi-universal C–Love relation. Column identifications and every headline independently re-derived; 5 contract tests pass. METHODOLOGICAL LESSON: in (M, R) the axes are incommensurable ($M_\odot$ vs km), so the degeneracy ANGLE ψ is metric-dependent and meaningless (joint-whitening forces orthogonal long axes — a spurious ψ -gap= 90° on every pair until caught); the scale-free shape invariant is the correlation $\rho(M, R)$. RESULT D1: the NICER pipeline/model systematic is VALUE-dominated — every Riley-vs-Miller and spot-model difference is a central-value offset at fixed orientation (orientation-only value fraction 0.90–1.00, $\Delta\rho \leq 0.08$); the sharpest is J0740, a +1.75 km radius offset (Miller 14.05 vs Riley 12.30) plus a $3.2\times$ larger Miller ellipse. This is the OPPOSITE of the cosmology analog (the KiDS statistic-dependence, E48, was a precision/shape artifact). RESULT D3: $\rho(M, R)$ is organized by MASS REGIME and robust across independent pipelines — low-mass J0030 $\rho \approx +0.86$ (steep M–R banana) vs high-mass J0740 $\rho \approx +0.25$ (flat top of the M–R curve), regime $\Delta\rho = 0.61$, consistent across pipelines (0.84/0.89/0.86; 0.29/0.21) — the NS analog of “the kernel sets orientation” (organizer = location on the mass–radius curve). D4 (exploratory, flagged): X-ray J0030 vs GW170817 at $\sim 1.4 M_\odot$ is MIXED — both a ~ 1.5 km value offset (GW prefers smaller R , $\sim 1.1\sigma$) and $\rho: 0.86 \rightarrow 0.28$, but the GW ρ is C–Love-mapping-induced. D5: all six posteriors are near-pure ($\text{Tr } \rho^2 = 0.990\text{--}0.996$); the $\sigma^2/8$ Chernoff repackaging (E62) persists. Honest limits carried up front: non-Gaussian (multimodal) NICER marginals, the mapping dependence of the GW point, and the QI repackaging.

GWTC-5/O4b: the curved chirp-mass law reconfirmed out-of-sample (E71)

E71 GWTC-5/O4b curved-law repeat. Third disjoint catalog after GWTC-1/2.1/3 (train) and GWTC-4/O4a (E67). Prereg locked (commit d19a7a8) BEFORE any O4b file was opened; decision rules byte-identical to E67, only a declared structural file/group adapter added. GWTC-5.0 O4b PE (Zenodo 20276106 + 20348006): 104 per-event releases, 52.84 GB, ZERO dropped, fully disjoint (event dates 240413–250119, zero name-overlap with the 86-event O4a set, 12 events from the 2025 tail). D1 (the on-record E65/E67 prediction; elongated $\text{axr} \geq 3$, $n = 32$): median $|\psi_{\text{curve}} - \psi_{\text{meas}}| = 1.22^\circ < 2^\circ$ — PASS (E67 gave 1.26°). D2: curve beats tangent, 5.41° vs 8.23° . D3: the E40 tangent law replicates a third time (median 8.23° ; chirp-vs-total wins 88%; Spearman(err, axr) = -0.69 , the strongest of the three catalogs). Zero free parameters — each prediction uses only the event’s own median source M_c and q marginal. Waveform-robust: the curve tracks the measured axis WITHIN every waveform family to $< 0.7^\circ$; the per-event largest-analysis pick (5 of 104 events do carry a Mixed group; the locked rule uses it) shifts only the absolute orientation. Verification four legs: 17/17 contract tests; independent numpy re-derivation of D1/D2; one event hand-recomputed from raw HDF5 via a different linear-algebra path (match to 0.004°); computed medians reproduce the release’s published PESummaryTable masses (exact for most, to $< 0.6\%$ for a few late high-mass events), column identity confirmed catalog-wide. The mass-plane geometry law is now reproducible across THREE independent catalogs spanning the full GW history. POST-HOC (labelled): the largest curve residuals are all round posteriors ($\text{axr} \sim 1$), the ill-defined-orientation regime the axr gate excludes — not law-breaking physics; the E72 outlier atlas must control for axr before reading residuals as anomalies.

E72 GW geometry outlier atlas. Prereg locked (0d1cebfb) BEFORE any residual-axis correlation. Where does the constant- M_c curve law break? Searched the 32 elongated O4b events ($\text{axr} \geq 3$; residual $r = \text{err}_{\text{curve}}$ from the locked E71 results) against 9 pre-declared axes with Benjamini–Hochberg FDR ($q = 0.05$) and partial-Spearman controlling axr; Tukey outlier fence; E59 coherence lens. RESULT — a clean PHYSICAL NULL: none of precession (χ_p), aligned spin (χ_{eff}), mass ratio, redshift, chirp/total mass, or secondary mass (NSBH/mass-gap proxy) survives FDR — the constant- M_c geometry has NO source-class-dependent breakdown in O4b (strengthens claim 1). The only surviving correlate is waveform-model disagreement (a systematic; $\rho = -0.64$, $q_{\text{BH}} = 0.001$, robust to drop-one ≤ 0.036), and its NEGATIVE sign says not-physics: the events the curve fits best ($r \sim 0.01$ –

0.2°) are those where waveforms disagree most — the curve tracks the waveform-ensemble consensus better than any single waveform tracks another (measurement optics; the E59 lesson again). One marginal Tukey outlier (GW240507, 3.37° vs fence 3.23°; low-SNR, asymmetric $q \sim 0.27$) logged as a WEAK preregisterable successor target, not a claim. Verify: 6 contract tests; the headline ρ and the Tukey fence independently re-derived (hand rank-Pearson; numpy percentile). Limits: $N = 32$ low power; loudness axis data-limited (6/32 picked groups carry a network-SNR column); eccentricity and higher-mode axes absent from the released PE (out of scope).

The loudest event: GW250114 no-hair deep-dive (E74)

E74 GW250114 deep-dive (loudest event). A single-event characterization + synthesis (NOT a locked battery): the O4b TGR products are not public, so the ringdown spectroscopy and no-hair verdict are the PUBLISHED result (Suzuki et al. 2026, arXiv:2605.03576, orthonormal QNMs; LVK IMR via GWTC-5). GW250114 is the loudest GW event to date (network SNR ~ 80), a near-EQUAL-mass BBH ($m_1 = 33.8$, $m_2 = 32.2 M_\odot$, $q \sim 0.96$, low spin), remnant $M_f^{\text{src}} = 62.8 M_\odot$, $a_f = 0.679$ — my PE extraction reproduces the published IMR remnant exactly (the data gate). INDEPENDENT contribution: propagating the remnant posterior through the Berti–Cardoso–Will Kerr fits gives the no-hair QNM targets — (2, 2, 0) at 248.9 Hz / $\tau = 4.10$ ms and the (2, 2, 1) overtone at 243.4 Hz / $\tau = 1.35$ ms (tight to ~ 1 Hz, the remnant being measured to $\sim 1.5\%$). The published analysis detects the (2, 2, 1) overtone at 99.9% (up from 82.5% non-orthogonal) and finds $\delta f_{221}, \delta \gamma_{221}$ consistent with zero — Kerr no-hair UPHELD at record SNR. GW250114 is essentially a louder GW150914-twin: near-identical remnant ($M_f \sim 68 M_\odot$ detector, $a_f \sim 0.68$; GW150914-like $f_{220} = 250$ Hz) but $\sim 3\times$ the SNR, which is why the once-contested overtone is now a clean detection. Curved-law N/A (near-equal-mass $\Rightarrow$ round posterior, $\text{axr} \sim 2$, outside the E71 elongated set). Verify: 4 QNM contract tests (GW150914 reference, $f \propto 1/M$, spin trend, overtone damps faster); published-value validation of all four remnant/mass quantities. Limit: the strain spectroscopy is not reproduced — the no-hair verdict is the cited published result; the area theorem awaits the O4b TGR release.

A geometric test of GR: the inspiral mass-combination exponent (E78)

E78 posterior-geometry test of GR (the inspiral exponent). EXPLORATORY/post-hoc, grown from E71/E72: since the (m_1, m_2) posterior orientation is a zero-parameter function of the constant- $\mathcal{M}_c$ curve, INVERT the law and measure which chirp-mass exponent the geometry prefers — a test of GR through posterior geometry, not waveform templates. Generalize the curve to $\mathcal{M}_p = (m_1 m_2)^p (m_1 + m_2)^{1-2p}$ (GR: $p = 3/5 = 0.600$, set by the 0PN phase $\propto \mathcal{M}_c^{-5/3}$); the orientation is SCALE-FREE ($m_1(q; p) \propto q^{-p}(1+q)^{2p-1}$), isolating p from the mass scale. On the 32 elongated O4b events: $\hat{p} = 0.616 \pm 0.011$ (stat) ± 0.004 (waveform) vs GR 0.600 — CONSISTENT at 1.3σ , a $\sim 2\%$ geometry-only confirmation of the inspiral mass-combination exponent. Validated: injection recovers the true exponent unbiased (bias $+0.0002$; $0.55 \rightarrow 0.550$, $0.65 \rightarrow 0.651$); the waveform systematic is negligible (± 0.004 — the per-event $\sim 3^\circ$ scatter is incoherent and averages out); 5 contract tests; sensitivity $d\psi/dp \sim 30^\circ/\text{unit}$. A deviation would have signalled modified inspiral phasing (dipole radiation / varying- G); none seen. HONEST: a methodologically independent cross-check, NOT more sensitive than the LVK -1PN dipole bound; O4b-only; the $\sim 1.2^\circ$ E71 law residual sets a per-event floor (~ 0.04 in p), so the small $+0.016$ offset is at the level of the law’s own imperfection, not claimed as physics. Successor (preregisterable): the fitted exponent is reproducible across GWTC-3/O4a/O4b and $= 0.600$.

E79 cross-catalog exponent — a 3σ “deviation” that is a systematic. Tests E78’s successor (prereg locked BEFORE O4a was opened): is the geometry-fit exponent reproducible across catalogs and $= 0.600$? On GWTC-4/O4a (86 events, disjoint from O4b): $\hat{p}(\text{O4a}) = 0.647 \pm 0.014$ (3.3σ from GR), $\hat{p}(\text{O4b}) = 0.616 \pm 0.011$; they REPRODUCE each other (D2 PASS) but combined $\hat{p} = 0.628 \pm 0.009$ sits a naive 3.1σ above 0.600 (D1 FAIL). The prereg pre-committed a follow-up before any claim — and it dissolves the anomaly: (i) the two independent catalogs disagree by $0.031 \approx$ the offset 0.028 , so the bootstrap stat error underestimates — a systematic floor ± 0.016 gives 1.5σ ; (ii) the offset is elongation-dependent (Spearman(p^* , axr) $= -0.41$, coherent in BOTH catalogs) and the cleanest, most-elongated half gives $\hat{p} = 0.606 \approx$ GR. By the coherence lens a cross-dataset-coherent anomaly is a SYSTEMATIC — here a finite-width / higher-order imperfection of the pure-0PN constant- $\mathcal{M}_c$ model, vanishing in the clean limit — NOT a GR violation. Honest headline: the effective exponent

is 0.628 ± 0.009 (stat) ± 0.016 (syst), GR-consistent at 1.5σ , cleanest events 0.606; the naive 3σ is a false alarm the discipline caught (the E47/E58 lesson, live). Verified: injection unbiased on both catalogs; disjointness = 0; column-ID clean; honest budget independently re-derived. Successor: a 1PN geometric model should reproduce the -0.4 elongation trend and flatten it to 0.6; add GWTC-3 as a third point.

The information anatomy of a detection (E73)

E73 information anatomy of a detection. CHARACTERIZATION (forward Fisher from the waveform model + PSD, NOT strain-extracted; the data-driven test is the successor). For all 190 real O4a+O4b events, the frequency by which 50%/90% of the Fisher information about each parameter is accumulated — chirp mass (0PN, $\sim f^{-5/3}$), mass ratio (1PN), spin (1.5PN). RESULT: the ordering $\mathcal{M}_c < \eta < \chi$ holds for ALL 190 events — chirp mass is always learned first (~ 35 Hz), spin last; the anatomy “richness” ($\chi_{f50} - \mathcal{M}_{c,f50}$) has Spearman = -1.00 with total mass (lighter / longer-inspiral $\Rightarrow$ wider anatomy; heavy events merge almost immediately in band and compress it). A frequency-resolved “when is each thing measured” anatomy of every detection. GR-test direction (a -1 PN dipole term): the RANK is reliable (Spearman $+1.00$ with mass; the lightest events — GW230529, the mass-gap merger, and GW230518 — are most sensitive, matching the known low-mass preference of dipole tests), but the absolute forecast is conditioning-limited (Fisher cond $\sim 10^{12}$; arbitrary β units), so only the ordering is quoted. Verify: 5 contract tests ($\mathcal{M}_c < \eta < \chi$ ordering, $f \propto 1/M$, $\sigma_\beta \propto 1/\rho$, lighter $\rightarrow$ tighter dipole); conditioning explicitly checked; ordering + richness independently re-derived; suite 37/37. Honest ceiling: model-PREDICTED, not strain-MEASURED — the data-driven test (re-analyse loud events in frequency slices, compare empirical vs GR-Fisher information accumulation) is the successor, along with a numerically-stable ppE dipole forecast versus the LVK -1 PN bound.

The GW inference manifold: the catalog as one geometric object (E82)

E82 the GW inference manifold. CHARACTERIZATION (post-hoc; radical idea made concrete). Treat all 190 O4a+O4b (m_1, m_2) posteriors as ONE object — each a Gaussian $N(\mu, \Sigma)$, the catalog the cloud of them, distance = 2-Wasserstein (location $|\Delta\mu|^2$ + Bures shape). RESULT: (1) location(population) 95% vs shape(measurement) 5% — mass-range-dominated, partly trivial; (2) the curved law slaves each ellipse’s orientation to its location (median 1.26° , E71; residual physics-blind, E72), so orientation is NOT an independent axis; (3) classical-MDS variance 97%/3%/0%, participation-ratio effective dimensionality = 1.07 (naive per-event DOF = 5); axis 1 = POPULATION ($|\rho_{\text{Mc}}| = 0.99$), thin axis 2 = PRECISION ($|\rho_{\text{axr}}| = 0.51$). VERDICT: the catalog of black-hole mergers, geometrically, is a population sequence (chirp mass) dressed with a thin, law-determined measurement structure — subtract the law and the population remains; measurement and population are geometrically separable, and here separated. The catalog-level form of the program’s thesis (E71 orientation = law(location); E72 residual physics-blind): the measurement optics occupy, and REMOVE, exactly one geometric axis. Verify: 4 contract tests (Bures closed form vs scipy `sqrtm` to 10^{-8} ; MDS recovers synthetic 1D/2D dimensionality); suite 41/41. Honest: the 95/5 split and eff-dim ~ 1 are mass-range-dominated (not oversold); Gaussian approximation; mass plane only. Successor: the full-parameter-space (spin, distance) Fisher-Rao version — the deferred E80 sloppy-spectrum program.

A strain-level pipeline, validated on GW250114 (E83)

E83 from-scratch strain pipeline. CAPABILITY + consistency check (“science first”): the strain-level analysis the exotic-test class requires, built with numpy/scipy only (no gwpy). Download GWOSC O4b strain $\rightarrow$ whiten (off-source Welch PSD) $\rightarrow$ bandpass $\rightarrow$ instantaneous frequency (Hilbert) for the chirp; post-peak dominant-frequency + log-envelope damping for the ringdown. GW250114 (H1): whitened peak 388σ (loudest to date); the CHIRP is recovered (f_{inst} rising through merger $89 \rightarrow 117 \rightarrow 187 \rightarrow 220$ Hz across ± 5 ms) and the RINGDOWN is recovered $f_{\text{dom}} = 241$ Hz (3.2% from E74’s Kerr 248.9 Hz), $\tau = 4.7$ ms (E74 4.10) — CONSISTENT with the E74 Kerr prediction. This independently cross-checks E74: the strain ringdown matches the Kerr QNM computed from the INSPIRAL-derived remnant — inspiral $\rightarrow$ remnant $\rightarrow$ ringdown closes in the

actual data. Verify: 4 contract tests (a 250 Hz/4 ms damped sinusoid injected into noise is recovered; whitening flattens red noise; the bandpass suppresses out-of-band tones $> 1000\times$; a chirp’s instantaneous frequency rises); suite 45/45. HONEST: a consistency check, NOT a measurement of record — single-detector, single-mode, fixed-start point estimate (few-% frequency / $\sim 15\%$ τ bias vs a proper Bayesian multi-mode ringdown). Successor: a Bayesian ringdown; the data-driven frequency-geometry test (E73 successor); the exotic echo / dark-matter-frequency-rotation tests ONLY after multi-event validation + coherence-lens/systematic controls (per E79/E72, a positive there is likelier a systematic than new physics).

E84 data-driven chirp mass — ballpark result, honest bias, feasibility boundary. First physics from the E83 pipeline: fit the textbook OPN law $f^{-8/3}(t) = K(t_c - t)$ to the spectrogram RIDGE of GW250114’s whitened H1 strain — a template-free chirp mass. Raw: $\mathcal{M}_c^{\text{det}} = 33.0 M_\odot$ vs PE 31.1 (+6%). BUT injection calibration shows the estimator is biased +15–30% HIGH (frequency quantization + window smearing in an accelerating sweep; grows with $\mathcal{M}_c$), so the raw agreement is PARTLY FORTUITOUS — honest claim: ballpark-consistent with a ~ 20 –30% template-free systematic, NOT precision. (The double-check caught this before it was banked — the E79 lesson again; a scipy STFT window-center double-offset bug was also caught by the contract tests.) NEGATIVE RESULT with mechanism: the band-resolved version (the frequency-geometry observable) is INFEASIBLE via ridge methods — GW250114 sweeps $64 \rightarrow 200$ Hz in < 25 ms, less than one STFT window (upper bands: 4 and 1 confident points), while light events are too faint per-slice; this is E73’s anatomy COMPRESSION observed in data. The data-driven frequency-geometry test therefore requires band-limited Bayesian PE (a major build), not more signal processing. Verify: 5 contract tests (slope-inversion roundtrip; recovery within the characterized bias envelope, never low; bias GROWS with $\mathcal{M}_c$; noise false-positive rate $< 15\%$); suite 50/50; the injection-bias table ships in the results JSON.

E85 Bayesian ringdown of GW250114 — RETRACTED, superseded by E87. The numbers in this entry are withdrawn: the posterior was PRIOR-DOMINATED (the whole 160–340 Hz grid sat within 4 lnL, so the “90% CI” was the grid) and was statistically indistinguishable from no signal. Two covariance defects cost $\sim 33\times$ in matched-filter SNR; see E87 for the diagnosis, the fix, and the validated measurement. The machinery below — analytic amplitude/phase marginalization, the Kerr inversion, and the truncation that defeats acausal whitening leakage — is sound and is reused unchanged. Retained verbatim as the record of the error. The Bayesian layer over the E83 pipeline: posterior on (f, τ) for the 220 mode $\rightarrow$ BCW inversion $\rightarrow$ a ringdown-ONLY (M_f, a_f) posterior vs the IMR remnant. Final likelihood: TRUNCATED time-domain (pyRing-style) — the window is compared to the model under the empirical noise AUTOCOVARANCE (Toeplitz + Cholesky), amplitude/phase marginalized analytically (linear params), (f, τ) gridded exactly (~ 1 s/fit, deterministic). RESULT (H1-only, 220-only): $f_{220} = 254$ [167, 333] Hz (start+3 ms) and 256 [169, 333] Hz (start+5 ms) — start-time STABLE, median at the IMR-Kerr prediction (248.6 Hz); ringdown-only remnant $M_f = 79$ [45, 122] $M_\odot$, $a_f = 0.87$ [0.19, 0.98] — the IMR remnant (68.4, 0.679) lies INSIDE the 90% region. A no-hair consistency check closed end-to-end with our own machinery, at the honest precision of a single-detector single-mode segment (ringdown SNR ~ 6). The build defeated three diagnosed artifacts, each now regression-guarded: (1) an amplitude prior that EXCLUDED the true $\sim 2 \times 10^{-21}$ strain scale; (2) an ensemble-sampler convergence failure (fixed by analytic linear-amplitude marginalization + exact gridding); (3) the important one — whiten-then-window ACAUSAL LEAKAGE: the 388σ merger rings forward through the PSD structure into the window, manufacturing start-time-unstable fake modes (192 vs 273 Hz); cured by the truncated autocovariance likelihood, and guarded by a regression test in which a $500\times$ -noise pre-window burst must yield a flat likelihood surface. Verify: 6 contract tests (sampler on a known Gaussian; BCW roundtrip to 10^{-6} ; grid injection recovery; surface peaks at injection; leakage regression); suite 56/56. Limits: H1-only, 220-only, fixed start, weak-consistency precision. Successor: L1 + coherent network, 221 overtone (4 linear amplitudes, same marginalization), start-time marginalization — then overtone/area-theorem questions in-house, and only then the exotic-echo class on this artifact-safe foundation.

E86 the band-resolved boundary, closed from both sides. E84 found band-resolved $\mathcal{M}_c$ infeasible on a HEAVY event and PREDICTED the complementary light-event failure; E86 tests that on the best case in the catalog and closes the method. Event: GW241011 — the MOST ELONGATED event (axr 17.3, curved-law residual 0.62°), $q = 0.32$, SNR ~ 36 , and a ~ 1.2 s inspiral ($7\times$ GW250114’s), H1+V1 public. RESULT: it fails here too. (1) Only the LOWEST band tracks a real chirp — ridge monotonicity Spearman(t, f) is +0.82 at 30–50 Hz but $-0.34, -0.02, -0.38$ at 40–70, 50–90, 80–150 Hz: line/noise structure, not a chirp (the conditioning already whitens AND notches a measured line forest — 60.00 Hz sits $\sim 400\times$ above the local median). That is

E84’s predicted per-slice-SNR failure. (2) The natural rescue — a LONG (0.5 s) window, affordable only because the low-frequency sweep is slow — buys monotonicity but is DEGENERATE: injections of known chirps into REAL off-source noise give recovered 7.60, 7.89, 4.88 for true 8.0, 10.9, 14.0, i.e. a 36% change in truth moves the estimate 4% — **response slope** 0.11 (1.0 unbiased, 0 no information). The real $\mathcal{M}_c^{\text{det}} = 8.35$ is “consistent” with the true ~ 10.9 only vacuously. VERDICT: ridge/spectrogram methods cannot deliver band-resolved $\mathcal{M}_c$ for ANY event class — heavy (sweep > window), light (faint per-slice + lines), long-window rescue (degenerate) — so the frequency-geometry observable REQUIRES matched-filter or banded Bayesian PE, now a measured requirement rather than a guess. Verify: 5 contract tests (notch suppresses a planted 60 Hz line; monotonicity discriminates injected chirp from noise; estimator $\sim$ unbiased where the window IS adequate; and the compressed response slope is itself regression-guarded at < 0.75); suite 61/61.

E87 E85 retracted — its ringdown posterior WAS the prior; covariance fixed, validated vs LVK pyRing. E85 reported $f_{220} = 253.9 [167, 333]$ Hz as an “honestly wide” Kerr-consistent measurement. It was not a measurement: the entire grid sat within **4 lnL**, so the interval was the grid, the IMR remnant fell inside because EVERYTHING did, and the “start-time stability” was a flat surface. THE TELL: adding L1 to H1 failed to tighten the posterior (162.0 vs 160.0 Hz) — impossible if there were information to add. Diagnosis: the fraction of the (f, τ) grid within 2 lnL of the peak was **1.00** on real data (0.02 on a loud injection), and both detectors railed at the grid’s LOW edge. ROOT CAUSE, two defects in the truncated-likelihood noise covariance, together $\sim 33\times$ of lost matched-filter SNR: (1) an ACF taper applied to force positive-definiteness crushed the off-diagonal correlations while keeping $\text{acf}[0]$, driving C toward white noise at the BROADBAND rms of high-passed strain, so a 250 Hz filter was charged 25–60 Hz seismic power ($8\times$); (2) a diagonal ridge of $10^{-3}\times$ the BROADBAND variance was itself a white floor far above the true in-band noise (a further $\sim 3.5\times$). FIX: build the ACF as the Wiener–Khinchin transform of a Welch PSD — positive-definite BY CONSTRUCTION, so no taper is needed and the noise colour survives — with the ridge at 10^{-6} . A bug introduced WHILE fixing it and caught by an explicit check: the one-sided PSD must be halved before the inverse FFT, or C comes out exactly $2\times$ too large, inflating every credible interval by $\sqrt{2}$ WITHOUT moving the peak (invisible to any peak-location test; caught only by testing $\text{acf}[0]$ against $\text{var}(\text{noise})$: ratio $2.0012 \rightarrow 1.0006$). VALIDATION GATE, the point of the exercise, PASSED: against LVK’s own pyRing Kerr₂₂₀ posteriors at 19 start times (Zenodo 17018009), mapped to f_{220} through the SAME BCW relation — median offset **2.0 Hz** over 2–18M (max 5.8), LVK’s median inside our 90% CI at **10/11** start times, and the credible-interval WIDTHS track (16.4 vs 16.3 Hz at 12M), i.e. the covariance is CALIBRATED, not merely centred. PHYSICS: the network start-time scan reproduces the early-time bias of a single-220 fit — 230 Hz with an **8 Hz** CI at t_{peak} (tight and WRONG), rising to a 251–254 Hz plateau at 12–18M_f, then back to the prior past 24M_f. *The tightest interval is the one that most confidently excludes the right answer.* Also withdrawn: a “ $251.7^{+5.1}_{-5.0}$ Hz LVK figure” used earlier in the session was never sourced and is not in the release; all comparison now reads LVK’s actual samples at runtime. Verify: 9 contract tests (WK normalization; PD-without-taper; corrected covariance $> 5\times$ the E85 construction; pure noise gives an UNINFORMATIVE posterior; inherited pre-window-burst leak guard; injection recovery; adding a detector must not broaden; Kerr roundtrip; peak location) — with the synthetic-noise fixture CALIBRATED (a gentle $60\times$ spectral step reproduces only a $1.8\times$ taper penalty vs $\sim 20\times$ for real H1, so it was steepened until it reproduced $\sim 40\times$; a fixture too tame to exhibit the bug would make the regression test vacuous — the same failure mode as the bug). Suite 70/70.

E88 the 221 overtone in GW250114 — a significant two-mode preference that is NOT the overtone. A negative with a named cause. GEOMETRY FIRST, computed before looking at any published result: at the IMR remnant the 221 sits 5.5 Hz from the 220 (243.1 vs 248.6) but decays $3\times$ faster (1.35 vs 4.10 ms), so over a 40 ms window the bases are $\sim 76\%$ **collinear** — most of the overtone is something the 220 can already absorb, and only the ORTHOGONALIZED $\text{SNR}_{\perp}$ is available: 6.9 at t_{peak} , 3.3 at 3M, 0.9 at 8M. So any overtone evidence must live at 0–3M and be gone by 6M; this independently predicts the profile of LVK’s own evidences. TWO METHOD RESULTS were needed first: (a) a Bayes factor needs a PROPER amplitude prior — E85/E87 marginalized amplitudes with a flat, improper one (fine for a posterior, undefined for model comparison); replaced by a Gaussian prior, formula verified against brute-force 2-D integration to 3×10^{-14} . (b) The naive multi-detector model DESTROYS the test: a free 221 amplitude per detector adds 4 linear parameters and costs ~ 11 **lnB of pure Occam penalty on signal-free data**. But the 221 and 220 share a detector’s antenna pattern, so the complex ratio $\rho = A_{221}/A_{220}$ is COMMON across detectors; writing $\text{Re}[A_{\text{det}}(T_{220} + \rho T_{221})]$ keeps it linear in A_{det} via $M_{\text{red}} = M_{\text{full}}K(\rho)$ (exact to 8×10^{-16}). Both models then carry two linear amplitudes per detector: σ_A cancels (at most 0.32, median 0.13, over a $33\times$ range), the whole overtone cost is the ρ prior volume (synthetic

injections with ρ contained: 2.00/2.12/1.83 vs $\ln 9 = 2.20$; real data 1.01–2.06, median 1.52, i.e. ρ is only weakly measured), and the signal-free penalty drops to +0.9. DOMINANT SYSTEMATIC, measured: with the remnant PINNED an injected $|\rho| = 1$ overtone gives $\Delta \ln B = +21.2$; with (M_f, a_f) MARGINALIZED it collapses to -0.05 — **essentially ALL absorbed** (at $|\rho| = 2$: $+110.0 \rightarrow +15.1$, 86%). The single-mode model just moves the remnant, peaking at (54.0, 0.150) whose 220 mode is $f = 234.8$ Hz, $\tau = 3.04$ ms. *That is E87’s early-time bias seen from the other side: same mechanism, two batteries.* Window length does not rescue it ($\Delta \ln B \sim 0$ –1 for 20–320 ms). RESULT: $\Delta \ln B = 6.15/6.09/6.33/5.88$ at 0/1/2/3*M*, still 4.94 at 8*M*, negative by 10*M*. Against the CORRECT null — a loud 220 with NO overtone injected into real noise, 13 realizations, -2.06 ± 0.98 — the observed +6.15 is **8.4 σ** , 0/13. (The signal-free null, $+0.91 \pm 0.58$, is the WRONG one: it contains no 220 either and would flatter the result to 9.4 σ .) THE DISCRIMINATOR: a genuine 221 must fade with a 2.01 ms e-folding, and an injected $|\rho_0| = 1.5$ template duly falls +5.69 from 0*M* to 8*M* (undetectable already by 1*M*); the observed preference falls only +1.21, persisting to 8*M* where $\text{SNR}_\perp = 0.94$ and no overtone can survive. **A preference that does not decay like an overtone is not the overtone** — by the standing coherence lens, a systematic of roughly constant size $\sim +5$. LVK’s profile DOES decay ($20.8 \rightarrow 10.1 \rightarrow 5.7 \rightarrow 0.7$ over 0–3*M*), matching the $\text{SNR}_\perp$ prediction; theirs, not ours, is the analysis that supports an overtone claim. We match them at 2, 8, 10 and 12*M* and differ most at 0–1*M*, where our looser detector model costs us. The systematic is NOT identified; sky-position-tied amplitudes are the first suspect. Verify: 10 contract tests; suite 80/80.

E89 spin magnitude rises with mass across 265 events — a hierarchical-assembly signature. THE PHYSICS: a comparable-mass BBH merger leaves a remnant with $a_f \approx 0.69$ almost independently of what went in — one of GR’s cleanest predictions, and visible directly in this catalog (**final_spin** sits at 0.690–0.707 in EVERY mass bin). Black holes that are themselves merger remnants should therefore carry $a \sim 0.7$ and sit preferentially at high mass. WHY E42–E44 SAW NOTHING: those batteries hunted this channel in χ_{eff} and found no requirement for it, which is exactly right — dynamical assembly gives ISOTROPIC orientations, so large spin magnitudes average to $\chi_{\text{eff}} \approx 0$. χ_{eff} is structurally blind to the population; the magnitude a_1 is not. This is E44’s question asked in the observable that can see it, on 265 events instead of 75. RESULT: hierarchical fit of $p(a_1|m_1) = \text{Beta}(\text{mean} = \text{sigmoid}(b_0 + b_1 \ln(m_1/30)))$ gives $b_1 = +1.193$, $\Delta \ln L = 10.94$, population mean a_1 running $0.083 \rightarrow 0.304 \rightarrow 0.499 \rightarrow 0.647$ at $m_1 = 8/30/60/100 M_\odot$ against the 0.69 remnant reference. Descriptive Spearman +0.554 ($p = 9 \times 10^{-23}$), repeating independently in GWTC-3 (+0.683), O4a (+0.490) and O4b (+0.475). The a_1 prior is uniform(0,1) — verified against the **priors** group — so each marginal posterior is proportional to its marginal likelihood. SIGNIFICANCE against nulls, not χ^2 , because heavier events are genuinely less informative (Spearman($\ln m_1$, concentration) = -0.517): forward simulation with mass-independent truth gives $b_1 = -0.046 \pm 0.138$; permuting masses within concentration strata gives -0.002 ± 0.173 . Both sit at ZERO, so the estimator is not biased by the confound, and the observed +1.193 exceeds the maximum of all **300** null trials. REMNANT-SPIN LEAK REFUTED: a leak would put heavy a_1 at **final_spin** (gap ~ 0 , per-event correlation $\sim +1$); the measured gap stays -0.197 above $60 M_\odot$ and the per-event correlation is ~ 0.4 – 0.5 for LIGHT events too. TWO TRAPS, both now contract-tested: (1) an “informative events” cut is CIRCULAR — among heavy events spin is well measured precisely WHEN IT IS LARGE (Spearman(a_1 , KS-from-prior) = +0.93 across the 10 heaviest), and the cut manufactures the very trend it appears to confirm; not used in the result. (2) The first null generator was WRONG and would have made this battery UNDER-report itself: modelling a posterior as a Beta centred on the truth makes a small true spin pile against $a_1 = 0$ and look informative, whereas real uninformative posteriors relax to the FLAT prior (real heavy low-information events sit at median $a_1 = 0.474$). The wrong shape produced a spurious null slope of $+0.42 \pm 0.11$ that would have been subtracted from +1.19 as estimator bias; fixed to a truth/uniform mixture. A tempting analytic argument — that a flat posterior contributes $\ln 1 = 0$ to every model and so cannot bias anything — holds only in the exact-flat limit; NEAR-flat posteriors do not obey it, which is why the nulls had to be run rather than reasoned about. LIMIT, and it is not small: **SELECTION EFFECTS ARE NOT MODELLED** — this is the DETECTED population, not the astrophysical one, and aligned spin raises SNR, so the astrophysical amplitude is not established without a VT/injection campaign. The effect is also not novel; what is new here is the independent measurement with the artifact and null calibration spelled out. Verify: 9 contract tests; suite 89/89. Successor: a 1G+2G mixture with a selection function, and the tilt-isotropy test.

E90 superradiance / ultralight bosons (SCOUT, not committed) — no exclusion survives calibration. Penrose’s rotational-energy extraction in its wave form: an ultralight boson forms a bound cloud around a spinning BH and spins it down, so if a boson of mass μ exists every BH must satisfy $a < a_{\text{crit}} = 4\alpha/(1 + 4\alpha^2)$ with $\alpha = GM\mu/\hbar c^3$ (derived from scratch; $\alpha = 0.7485$ for $M = 10 M_\odot$, $\mu = 10^{-11}$ eV vs the literature’s ~ 0.75).

Applied to the 265 E89 posteriors, a naive analysis “excludes” $\mu \in [2 \times 10^{-14}, 4 \times 10^{-12}]$ eV. FOUR successive controls each killed a version: (i) $P(a > a_{\text{crit}})$ is prior-dominated (the uniform spin prior alone gives 0.5–0.95); (ii) a Bayes factor against a UNIFORM null makes the boson look FAVOURED ($\ln B = +40$), but shuffling spins across events also gives +16 — real spins are simply low, so ANY truncation wins; (iii) using the E89 astrophysical population as the null (the right structure) gives $\ln B = -1403$, but $-690 = \ln 10^{-300}$ is the numerical floor: single events’ posterior TAILS, exactly where PE is least trustworthy; (iv) with an E79-style systematics floor ($\epsilon = 0.1$, capping any one event at $\ln 0.1$), $\ln B = -69$ — but mock catalogs drawn from the E89 population with *no boson at all* give -63.0 ± 2.8 . VERDICT: the real data sit **1.95 σ** below the no-boson expectation; the “excluded band” was a look-elsewhere artifact across 60 μ values. **No boson mass is excluded.** Kept as a scout (no src/results committed): the reusable content is the $\alpha/a_{\text{crit}}/\tau_{\text{SR}}$ machinery and the demonstration that E89’s population is the correct null for exotic-spin tests.